

COMS4040A & COMS7045A Assignment 3

Hand-out date: 10:00 am, April 28, 2020

Due date: 24:00 pm, May 28, 2020

Instructions

1. This is an individual assignment. Adhere to the academic integrity. Plagiarism will result in 0%.
2. The due date is strictly before 24:00 pm, May 28, 2020.
3. Before the final hand in, submit your write-up PDF file to Turnitin to produce a similarity check report. This report will be viewed by the lecturer only.
4. Hand in the electronic files and source codes on Sakai course site. See more instructions in Section **Hand-in**.
5. The total marks available for this assignment is 60 for both Hons students and MSc-CWRR students.

Outcome

1. Write programs in CUDA for CPU + GPU systems;
2. Optimize the performance of a CUDA program by using CUDA memory hierarchies efficiently.

1 Image Convolution Using CUDA C

1.1 Introduction

Convolution is an array operation where each output data element is a weighted sum of a collection of neighbouring input elements. The weights used in the weighted sum calculation are defined by an input mask array, referred to as the convolution mask here. The same convolution mask is typically used for all elements of the array. Convolution is commonly used in various forms in signal processing, digital recording, image processing, video processing, and computer vision. In these application areas, convolution is often performed as a filter that transforms signals and pixels into more desirable values.

Convolution typically involves a significant number of arithmetic operations on each data element. For large data sets such as high-definition images and videos, the amount of computation can be very large. Each output data element can be calculated independently of each other, a desirable trait for parallel computing. On the other hand, there is substantial level of input data sharing among output data elements with somewhat challenging boundary conditions. This makes convolution an important use case of sophisticated tiling methods and input data staging methods.

When applied in image processing tasks, convolution leads to results such as noise removal, edge detection and sharpening of details etc. If an image is represented as a 2D discrete signal $Y \in \mathbb{Z}^{M \times N}$, we can perform the (discrete) convolution




Figure 1: Performing convolution operation

in 2-dimension using a kernel $F \in \mathbb{Z}^{L \times P}$ as

$$(Y * F)[i, j] = \sum_{l=0}^{L-1} \sum_{p=0}^{P-1} Y[i-l, j-p] F[l, p], \quad 0 \leq i < M, \quad 0 \leq j < N. \quad (1)$$

The convolution operation in Equation (1) is actually a scalar product of the filter weights and all pixels of the image within a window that is defined by the extent of the filter and a centre pixel. Figure 1 illustrates the convolution using a small 3×3 kernel. The design of the convolution filter requires careful selection of the weights to achieve desired results.

1.2 Implementation Considerations

Image convolution can be efficiently implemented on massively parallel hardware, since the same operator gets executed independently for each image pixel. A naive implementation would simply execute the convolution for each pixel by one CUDA thread, read all values inside the filter area from the global memory, and write the result. (Note that this approach is inefficient.)

In this assignment, you are going to implement image convolution using the following methods.

1. Serial computation [8]
2. CUDA implementation using both global memory and constant memory [12]
3. CUDA implementation using both shared memory and constant memory [16]
4. CUDA implementation using texture memory [12]

In your implementation:

1. The filters in Table 1 can be used for testing the convolution result.
2. A pgm test image is given. To load a pgm image or write a pgm image, you may use the relevant CUDA SDK functions. A CUDA SDK sample program is provided in this regard.

3. The filtered values outside the image boundaries, referred to as ‘ghost cells’, should be treated as 0 values.
4. Your code should be written in a way that the size of convolution mask is arbitrary where applicable, that is, you should test it on different sizes of convolution masks, e.g., 3 by 3, 5 by 5, 7 by 7 etc.

Sharpening	Edge Detection	Averaging
$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

Table 1: Convolution masks

2 Hand-in

1. Your submission should be a single compressed file named as `yourStudentNo_hw3.tar.gz` that when extracts, gives me a folder named `yourStudentNo_hw3`. In this folder, you may include subfolders of source files for different approaches, `Makefiles` to build your code, a `readme` file on building and running your program, as well as a report named `report.pdf`. This document will be the main source of evaluation of the quality of your work.
2. Here is the list of contents and discussions you should put in your report.
 - A description of the design choices you tried and how did they affect the performance.
 - Performance comparison among all approaches implemented.
 - Performances using different sizes of input images, and using different sizes of ‘averaging’ convolution masks. (You don’t need to change the size of other types of convolution masks.)
 - Display output images from each implementation using a format with 2-4 images in a row, with clear captions and proper references in the text.
 - Include a “Questions and answers” section, where give short answers to the following questions:
 - (a) How many floating point operations are performed in your convolution kernel using global memory? Explain. [2]
 - (b) How many global memory reads are performed by your kernel using global memory and kernel using shared memory, respectively? [2]
 - (c) How many global memory writes are performed by your convolution kernel using shared memory ? [1]
 - (d) What would happen to the performance of your kernel using shared and constant memory when the size of the ‘averaging’ mask increases (say, to a substantial large size) ? [2]
 - A discussion on using different types of CUDA device memory.
 - References, including open source code projects, should be cited properly in your report.
 - Your report should be in single column format, use 11pt, single line spacing, and if any tables or diagrams are in it, they should be clearly readable. Furthermore, diagrams, figures, and other floating objects should be properly sized to shorten the length of your report. Finally, write in a style of scientific writing, and please avoid putting in non-technical stuff in your report. [5]